

Deep Learning Techniques

DL 3. Convolutional Neural Networks (CNN)

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

1

1

Outline

- Motivation
- CNN Operations
- Gradient-based Optimization
- LM with word-level and character-level CNNs

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

2

2

Comparison with RNN-based models

- RNN
 - Long-term dependency
 - Sequential computation (slow), prone to gradient vanishing
- CNN
 - N-gram-like models focusing on local dependency
 - Computation is faster (easier to be parallelized)
- **RNN and CNN can be used in combination!**

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

3

3

What is Convolution?

Dictionary

Search for a word



con·vo·lu·tion

/ˌkɑnvəˈlooʃən/

noun

1. a thing that is complex and difficult to follow.
"the convolutions of farm policy"
synonyms: complexity, intricacy, complication, twist, turn, entanglement, contortion; [More](#)
2. a coil or twist, especially one of many.
"crosses adorned with elaborate convolutions"
synonyms: twist, turn, coil, spiral, twirl, curl, helix, whorl, loop, curlicue, kink, sinuosity; [More](#)

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

4

4

What is Convolution?

- Wikipedia <https://en.wikipedia.org/wiki/Convolution>
 - **convolution** is a mathematical operation on two functions (f and g) to produce a third function that expresses how the shape of one is modified by the other.

$$(f * g)(t) \triangleq \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau.$$

An equivalent definition is (see [commutativity](#)):

$$(f * g)(t) \triangleq \int_{-\infty}^{\infty} f(t - \tau)g(\tau) d\tau.$$

- Essentially, it is an operation over f and g by taking a weighted sum of f (or g) using g (or f) as the weights.

Convolutional Neural Networks (CNNs)

- A type of neural networks
- Popular in image/video recognition, text mining, recommendation, etc.
- Inspired by biological processes
 - The *receptive (activated) field* of cortical neurons can be approximated mathematically by a *convolution operation*.

Convolution Operations over an Image (cont'd)

AI Shack <http://aishack.in/tutorials/image-convolution-examples/>

Blur Filter
(local average)

| | | |
|-----|-----|-----|
| 1/9 | 1/9 | 1/9 |
| 1/9 | 1/9 | 1/9 |
| 1/9 | 1/9 | 1/9 |



1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

9

| | | |
|----|----|----|
| -1 | -1 | -1 |
| 2 | 2 | 2 |
| -1 | -1 | -1 |

Horizontal
Line Filter



9

Convolution Operations over an Image (cont'd)

AI Shack <http://aishack.in/tutorials/image-convolution-examples/>

| | | |
|----|----|----|
| -1 | -1 | -1 |
| -1 | 8 | -1 |
| -1 | -1 | -1 |

Edge Filter

Below result I got with edge detection:



1/25/2024

10

10

Terminology & Notation

- **Input X** (or function f in our introduction of convolution)
 - Numerical representation of an image, a sentence, a time series, etc.
- **Filter (kernel) W** (or function g in our introduction of convolution)
- **Receptive Field**: the local region that the filter is applied to.
- **Feature Map**: the output of convolution

CNN Operations

- ✓ Convolution
- Striding
- Padding
- Pooling

Image Convolution (a toy example)

Animation: http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution
 More references [here](#)

| | | | | |
|---|---|---|---|---|
| 1 <small>$\times 1$</small> | 1 <small>$\times 0$</small> | 1 <small>$\times 1$</small> | 0 | 0 |
| 0 <small>$\times 0$</small> | 1 <small>$\times 1$</small> | 1 <small>$\times 0$</small> | 1 | 0 |
| 0 <small>$\times 1$</small> | 0 <small>$\times 0$</small> | 1 <small>$\times 1$</small> | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |

Convolved
Feature Map

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

13

13

Learnable Parameters in the CNN filter

- Filter size is $k \times k$ (learnable model parameters)
 - Given an $n \times n$ input, the produced feature map is $m \times m$ where $m = n - k + 1$.
- Compared to a fully connected network (perceptron) for a $R^{n \times n} \rightarrow R^{m \times m}$ linear mapping, its number of model parameter is $n^2 \times m^2$ (much larger than k^2).
- Therefore, CNN is **computationally more feasible** for extracting lower dimensional features from input data and is **less prone to overfitting**.

1/25/2024


@Yiming Yang, MLwG S24 DL3 CNN

14


14

Striding Operation in CNN

- striding size (hyperparameter): the number of units to shift step-by-step



Striding of size 1 ("no striding")



Striding of size 2

Animation: https://github.com/vdumoulin/conv_arithmetic

1/25/2024


@Yiming Yang, MLwG S24 DL3 CNN

15


15

Padding in CNN

- Adding zeros around the input image (for desirable output size, or focusing on edges)



No padding



Padding

Source: https://github.com/vdumoulin/conv_arithmetic

1/25/2024


@Yiming Yang, MLwG S24 DL3 CNN

16

16

Pooling in CNN

- Fixed operation (average or max) over each local region for further dimensionality reduction



Max pooling with 2x2 Filter.

Source: <http://cs231n.github.io/convolutional-networks/#pool>


1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

17

17

Image Pooling: Capturing the Local Invariance



Max pooling with 2x2 Filter.

Source: <http://cs231n.github.io/convolutional-networks/#pool>

Comparing the images on the left and the right: The local pattern (orange box) is shifted upwards by 1 unit on the right, but the output of pooling remain unchanged.

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

18

18

CNN for image object detection




Figure 3: A simple ConvNet. Source [5]

- Automatically learn the filters (kernels) based on labeled training data
- Extracting **local** and **lower-dimensional** features from input data
- Computationally more efficient than MLPs or RNNs
- **Why do we need multiple filters in parallel?**

1/25/2024

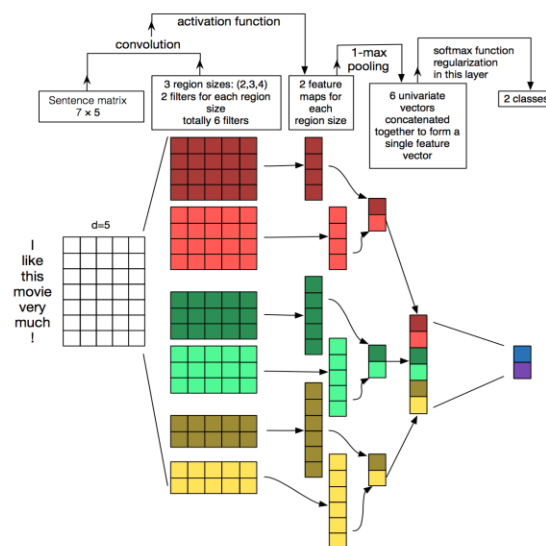
@Yiming Yang, MLwG S24 DL3 CNN

19

19

CNNs for Text

- Use word embedding to obtain the input “image” (1-D)
- Set the sizes of convolution filters to be $m \times d$ for $m = 2, 3, 4, \dots$
 - m is the number of words a filter takes into account (usually 1-5, like n-gram)
 - d is the size of word embedding
- Use multiple filters for each size (why?)



Typical CNN structure for text classification [Kim 2014]

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

20

20

Regular Convolutions with Stacked Layers

<https://stats.stackexchange.com/questions/287774/wavenet-is-not-really-a-dilated-convolution-is-it>

Non dilated Causal Convolutions



- This is a convolution with a filter size of 2 and a stride of 1, repeated for 4 layers.
- The receptive-field size (n) grows linearly in the numbers of the layers.


1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

21

21

Dilated CNN



- Skipping input values with certain steps at each layer (e.g., skipping 1, 3, 5, ... in Dilation 1)
- The receptive-field size grows exponentially in the number of layers ($\log n$).

1/25/2024


@Yiming Yang, MLwG S24 DL3 CNN

22

22

Animation of Dilated CNN

<https://stats.stackexchange.com/questions/287774/wavenet-is-not-really-a-dilated-convolution-is-it>



For the same receptive-filter size, **dilated CNN is more compact** than conventionally stacked CNN.

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

23

23

Outline

- ✓ Motivation
- ✓ CNN Operations
- **Gradient-based Optimization**
- LM with word-level and character-level CNNs

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

24

24

Gradient-based Optimization in CNN

- Denote the loss function as $L(y, \hat{y}_W(x))$
 - \hat{y}_W depends on the steps of (padded) convolution and pooling;
 - W consists of the matrix of the learnable filter weights.
- Model Training
 - Optimizing W through backpropagation
 - Differentiable $\frac{\partial \hat{y}_W(x)}{\partial W}$ in both the convolution and pooling steps
 - Leveraging parallel computing on GPU

1/25/2024


@Yiming Yang, MLwG S24 DL3 CNN

25

25

Optimization in CNN: A Toy Example (1D output)

- $X \in \mathbb{R}^{5 \times 5}$: the input matrix
- $W \in \mathbb{R}^{3 \times 3}$: the kernel matrix
 - Indexed by (k, l)
- $H \in \mathbb{R}^{3 \times 3}$: the feature map;
 - Indexed by (i, j)
- $\hat{Y} \in \mathbb{R}$: the system's output
- $Y \in \mathbb{R}$: the ground truth



1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

26

26

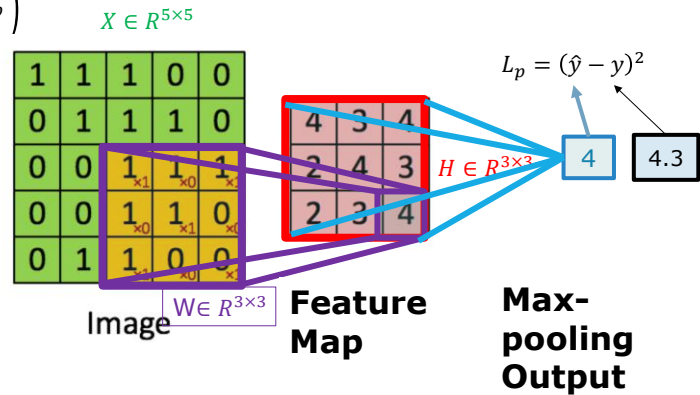
Optimization via stochastic gradient descent

$$W_{mn}^{(new)} := W_{mn}^{(old)} - \eta_k \nabla_{W_{mn}} \left(\frac{1}{|B|} \sum_{p \in B} L_p \right)$$

$$:= W_{mn}^{(old)} - \eta_k \frac{1}{|B|} \sum_{p \in B} \frac{\partial L_p}{\partial W_{mn}}$$

- $m, n \in \{1, 2, 3\}$
- L_p is the loss function on a training pair
- B is a mini-batch of training pairs
- η_k : learning rate at step k

$$\frac{\partial L_p}{\partial W_{mn}} = \frac{\partial L_p}{\partial \hat{y}} \sum_{i,j} \frac{\partial \hat{y}}{\partial H_{ij}} \frac{\partial H_{ij}}{\partial W_{mn}}$$



1/25/2024


@Yiming Yang, MLwG S24 DL3 CNN

27

27

Backpropagation for Max-pooling

- $\frac{\partial L_p}{\partial \hat{y}}$ is trivial, $\frac{\partial L_p}{\partial \hat{y}} = 2(\hat{y} - y)$
 - $\frac{\partial \hat{y}}{\partial H_{11}} := 1$
 - $\frac{\partial \hat{y}}{\partial H_{12}} := 0$
 - $\frac{\partial \hat{y}}{\partial H_{21}} := 0$
 - $\frac{\partial \hat{y}}{\partial H_{22}} := 1$
 - ...
- the index of the maximum in H



1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

28

28

Backpropagation for Convolution

$$\frac{\partial \hat{y}}{\partial W_{mn}} := \sum_{i,j \in \{1,2,3\}} \frac{\partial \hat{y}}{\partial H_{ij}} \frac{\partial H_{ij}}{\partial W_{mn}} := \sum_{i,j \in \{1,2,3\}} \frac{\partial \hat{y}}{\partial H_{ij}} X_{i+m-1,j+n-1}$$

- This acts just like another filter (convolution operation) for a weighted sum of the input cells.
- You can verify $\frac{\partial H_{ij}}{\partial W_{mn}} = X_{i+m-1,j+n-1}$ as following

$$H_{11} = W_{11}X_{11} + W_{12}X_{12} + W_{13}X_{13}, \quad H_{12} = W_{11}X_{12} + W_{12}X_{13} + W_{13}X_{14}, \quad H_{13} = W_{11}X_{13} + W_{12}X_{14} + W_{13}X_{15}$$

$$H_{21} = W_{11}X_{21} + W_{12}X_{22} + W_{13}X_{23}, \quad H_{22} = W_{11}X_{22} + W_{12}X_{23} + W_{13}X_{24}, \quad H_{23} = W_{11}X_{23} + W_{12}X_{24} + W_{13}X_{25}$$

$$\frac{\partial H_{11}}{\partial W_{11}} = X_{11}, \quad \frac{\partial H_{11}}{\partial W_{12}} = X_{12}, \quad \frac{\partial H_{11}}{\partial W_{13}} = X_{13}, \quad \frac{\partial H_{21}}{\partial W_{11}} = X_{21}, \quad \frac{\partial H_{22}}{\partial W_{12}} = X_{23}, \quad \frac{\partial H_{23}}{\partial W_{13}} = X_{25} \quad \dots$$

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

29

29

Text classification results in error (smaller is better)

[Xiang Zhang, Junbo Zhao, Yann LeCun, NIPS 2015]

Table 4: Testing errors of all the models. Numbers are in percentage. "Lg" stands for "large" and "Sm" stands for "small". "w2v" is an abbreviation for "word2vec", and "Lk" for "lookup table". "Th" stands for thesaurus. ConvNets labeled "Full" are those that distinguish between lower and upper letters

Non-neural models

| Model | AG | Sogou | DBP | Yelp P. | Yelp F. | Yah. A. | Amz. F. | Amz. P. |
|--------------------|-------|-------|------|---------|---------|---------|---------|---------|
| BoW | 11.19 | 7.15 | 3.39 | 7.76 | 42.01 | 31.11 | 45.36 | 9.60 |
| BoW TFIDF | 10.36 | 6.55 | 2.63 | 6.34 | 40.14 | 28.96 | 44.74 | 9.00 |
| ngrams | 7.96 | 2.92 | 1.37 | 4.36 | 43.74 | 31.53 | 45.73 | 7.98 |
| ngrams TFIDF | 7.64 | 2.81 | 1.31 | 4.56 | 45.20 | 31.49 | 47.56 | 8.46 |
| Bag-of-words | 16.91 | 10.79 | 9.55 | 12.67 | 47.46 | 39.45 | 55.87 | 18.39 |
| LSTM | 13.94 | 4.82 | 1.45 | 5.26 | 41.83 | 29.16 | 40.57 | 6.10 |
| Lg. w2v Conv. | 9.92 | 4.39 | 1.42 | 4.60 | 40.16 | 31.97 | 44.40 | 5.88 |
| Sm. w2v Conv. | 11.35 | 4.54 | 1.71 | 5.56 | 42.13 | 31.50 | 42.59 | 6.00 |
| Lg. w2v Conv. Th. | 9.91 | - | 1.37 | 4.63 | 39.58 | 31.23 | 43.75 | 5.80 |
| Sm. w2v Conv. Th. | 10.88 | - | 1.53 | 5.36 | 41.09 | 29.86 | 42.50 | 5.63 |
| Lg. Lk. Conv. | 8.55 | 4.95 | 1.72 | 4.89 | 40.52 | 29.06 | 45.95 | 5.84 |
| Sm. Lk. Conv. | 10.87 | 4.93 | 1.85 | 5.54 | 41.41 | 30.02 | 43.66 | 5.85 |
| Lg. Lk. Conv. Th. | 8.93 | - | 1.58 | 5.03 | 40.52 | 28.84 | 42.39 | 5.52 |
| Sm. Lk. Conv. Th. | 9.12 | - | 1.77 | 5.37 | 41.17 | 28.92 | 43.19 | 5.51 |
| Lg. Full Conv. | 9.85 | 8.80 | 1.66 | 5.25 | 38.40 | 29.90 | 40.89 | 5.78 |
| Sm. Full Conv. | 11.59 | 8.95 | 1.89 | 5.67 | 38.82 | 30.01 | 40.88 | 5.78 |
| Lg. Full Conv. Th. | 9.51 | - | 1.55 | 4.88 | 38.04 | 29.58 | 40.54 | 5.51 |
| Sm. Full Conv. Th. | 10.89 | - | 1.69 | 5.42 | 37.95 | 29.90 | 40.53 | 5.66 |
| Lg. Conv. | 12.82 | 4.88 | 1.73 | 5.89 | 39.62 | 29.55 | 41.31 | 5.51 |
| Sm. Conv. | 15.65 | 8.65 | 1.98 | 6.53 | 40.84 | 29.84 | 40.53 | 5.50 |
| Lg. Conv. Th. | 13.39 | - | 1.60 | 5.82 | 39.30 | 28.80 | 40.45 | 4.93 |
| Sm. Conv. Th. | 14.80 | - | 1.85 | 6.49 | 40.16 | 29.84 | 40.43 | 5.67 |

Red: worst performance on each dataset

Parameters/hyperparameters have a large influence!

Blue: best performance on each dataset

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

30

30

CNN vs. non-neural methods

- Non-neural methods (old classifiers)
 - Bag-of-words (BOW) features of input text (with binary or TF-IDF term weighting) assumes **independency** among words.
 - Bag-of-nGram features of input text can capture **local dependencies** with a significantly enlarged vocabulary (and high dimensional input vectors).
- CNN Models
 - Convolution can capture **local dependencies** with low-dimensional vectors in an efficient way.

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

31

31

Parameters and Other Factors

- Hyperparameter **matters!**
- Hyperparameters include
 - Convolution/pooling size
 - Striding/padding size
 - Dilation scope
- Optimization Algorithms
 - SGD/Ada-grad
 - Mini-batch size in SGE
- In text classification
 - Embedding methods

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

32

32

Embedding Strategies

- Pre-training vs. No pre-training
 - If an unlabeled large corpus is available, pretraining of word embedding can be used for faster fine-tuning and better generalization ability
- Word-level/Character-level embedding
 - Word embedding can be initialized with pretrained embedding, missing words are randomly initiated
 - If too many words are unknown in pretraining, character-level embedding may be used

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

33

33

Outline

- ✓ Motivation
- ✓ CNN Operations
- ✓ Gradient-based Optimization
- LM with word-level and character-level CNNs

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

34


34

Language modeling with character-level CNN

[Yoon Kim et al. AAAI 2016]

- Input
 - character-level embedding
- Output
 - Word-level embedding
- Aimed Advantage
 - Capture OOV (Out-Of-Vocabulary) words

CNN-generated embedding for "absurdity"



1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

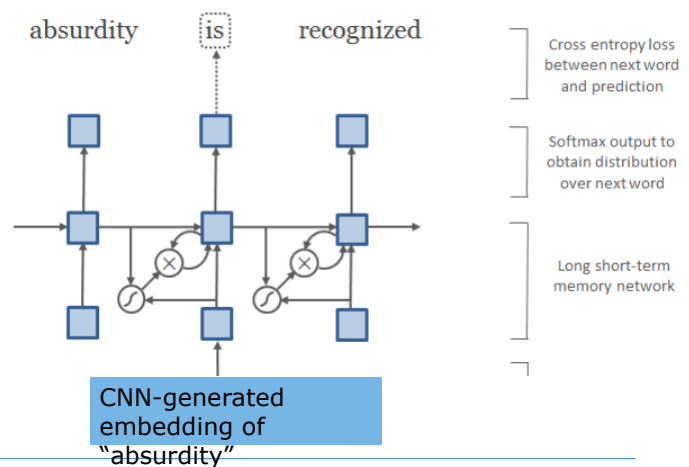
moment the absurdity is recognized

35

Language Modelling with CNN+LSTM

[Yoon Kim et al. AAAI 2016]

- The CNN-generated embedding is then fed into a LSTM for predicting the current word given the past context.



1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

36

36

CNN-generated embedding

[Yoon Kim et al. AAAI 2016]

| | In Vocabulary | | | | | Out-of-Vocabulary | | |
|-------------------------------|---------------------|--------------|----------------------|-----------------|--------------------|------------------------|---------------------|------------------|
| | <i>while</i> | <i>his</i> | <i>you</i> | <i>richard</i> | <i>trading</i> | <i>computer-aided</i> | <i>misinformed</i> | <i>loooooook</i> |
| LSTM-Word | <i>although</i> | <i>your</i> | <i>conservatives</i> | <i>jonathan</i> | <i>advertised</i> | – | – | – |
| | <i>letting</i> | <i>her</i> | <i>we</i> | <i>robert</i> | <i>advertising</i> | – | – | – |
| | <i>though</i> | <i>my</i> | <i>guys</i> | <i>neil</i> | <i>turnover</i> | – | – | – |
| | <i>minute</i> | <i>their</i> | <i>i</i> | <i>nancy</i> | <i>turnover</i> | – | – | – |
| LSTM-Char (before highway) | <i>chile</i> | <i>this</i> | <i>your</i> | <i>hard</i> | <i>heading</i> | <i>computer-guided</i> | <i>informed</i> | <i>look</i> |
| | <i>whole</i> | <i>hhs</i> | <i>young</i> | <i>rich</i> | <i>training</i> | <i>computerized</i> | <i>performed</i> | <i>cook</i> |
| | <i>meanwhile</i> | <i>is</i> | <i>four</i> | <i>richer</i> | <i>reading</i> | <i>disk-drive</i> | <i>transformed</i> | <i>looks</i> |
| | <i>white</i> | <i>has</i> | <i>youth</i> | <i>richter</i> | <i>leading</i> | <i>computer</i> | <i>inform</i> | <i>shook</i> |
| LSTM-Char (after highway) | <i>meanwhile</i> | <i>hhs</i> | <i>we</i> | <i>eduard</i> | <i>trade</i> | <i>computer-guided</i> | <i>informed</i> | <i>look</i> |
| | <i>whole</i> | <i>this</i> | <i>your</i> | <i>gerard</i> | <i>training</i> | <i>computer-driven</i> | <i>performed</i> | <i>looks</i> |
| | <i>though</i> | <i>their</i> | <i>doug</i> | <i>edward</i> | <i>traded</i> | <i>computerized</i> | <i>outperformed</i> | <i>looked</i> |
| | <i>nevertheless</i> | <i>your</i> | <i>i</i> | <i>carl</i> | <i>trader</i> | <i>computer</i> | <i>transformed</i> | <i>looking</i> |

Table 6: Nearest neighbor words (based on cosine similarity) of word representations from the large word-level and character-level (before and after highway layers) models trained on the PTB. Last three words are OOV words, and therefore they do not have representations in the word-level model.

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

37

37

Performance

[Yoon Kim et al. AAAI 2016]

| | <i>PPL</i> | <i>Size</i> |
|---|------------|-------------|
| LSTM-Word-Small | 97.6 | 5 m |
| LSTM-Char-Small | 92.3 | 5 m |
| LSTM-Word-Large | 85.4 | 20 m |
| LSTM-Char-Large | 78.9 | 19 m |
| KN-5 (Mikolov et al. 2012) | 141.2 | 2 m |
| RNN [†] (Mikolov et al. 2012) | 124.7 | 6 m |
| RNN-LDA [†] (Mikolov et al. 2012) | 113.7 | 7 m |
| genCNN [†] (Wang et al. 2015) | 116.4 | 8 m |
| FOFE-FNNLM [†] (Zhang et al. 2015) | 108.0 | 6 m |
| Deep RNN (Pascanu et al. 2013) | 107.5 | 6 m |
| Sum-Prod Net [†] (Cheng et al. 2014) | 100.0 | 5 m |
| LSTM-1 [†] (Zaremba et al. 2014) | 82.7 | 20 m |
| LSTM-2 [†] (Zaremba et al. 2014) | 78.4 | 52 m |

Table 3: Performance of our model versus other neural language models on the English Penn Treebank test set. *PPL* refers to perplexity (lower is better) and size refers to the approximate number of parameters in the model. KN-5 is a Kneser-Ney 5-gram language model which serves as a non-neural baseline. [†]For these models the authors did not explicitly state the number of parameters, and hence sizes shown here are estimates based on our understanding of their papers or private correspondence with the respective authors.

1/25/2024

38

38

Comparison with RNN-based models

- RNN
 - Long-term dependency
 - Sequential computation (slow), prone to gradient vanishing
- CNN
 - N-gram-like models focusing on local dependency
 - Computation is faster (easier to be parallelized)
- **RNN and CNN can be used in combination!**

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

39

39

References

- An Intuitive Explanation of Convolutional Neural Networks, ujjwalkarn, 2016
<https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>
- AI Shack <http://aishack.in/tutorials/image-convolution-examples/>
- Character-Aware Neural Language Models. Yoon Kim, Yacine Jernite, David Sontag, Alexander M. Rush., AAAI 2016
- Character-level Convolutional Networks for Text Classification, Xiang Zhang, Junbo Zhao, Yann LeCun. NIPS 2015. <http://papers.nips.cc/paper/5782-character-level-convolutional-networks-for-text-classification.pdf>

1/25/2024

@Yiming Yang, MLwG S24 DL3 CNN

40

40