

Neural Network Solvers for Combinatorial Optimization

Graph 10. Non-autoregressive CO Solvers

1

Lectures on Neural-network CO Solvers

- Graph 9: Autoregressive (AR) CO Solvers
- Graph 10: Non-autoregressive (NAR) CO Solvers
 - Reinforcement learning with DIMES [R Qiu*, Z Sun* & Y Yang, **NearIPS 2022**]
 - Supervised learning with DIFUSCO [Z Sun & Y Yang, **NearIPS 2023**]
- Graph 11: Pre-trained Large Language Models for CO
- Graph 12: Neural Solvers for Mixed Integer Programming

2

DIMES: A differentiable meta solver for CO

[Qiu*, Sun* & Yang, NearIPS 2022]

Key Ideas

- 1) A unified framework for CO problems (including TSP and MIS)
- 2) A **NAR** CO solver with an **encoder-only architecture** (removing the expensive decoding parts in the AR solvers), scaled to graphs with $n = 10,000$ nodes
- 3) Using **meta learning** (in the training-phase) to enhance the accuracy on new graphs

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

3


3

Unified Framework for CO Problems

- Convert any CO problem into a task of $\{0,1\}$ -vector generation
 - Let $\mathcal{F}_s = \{0,1\}^N$ be the space of **feasible solutions** for instance graph s ;
 - Let $c_s: \mathcal{F}_s \rightarrow \mathbb{R}$ be the **cost function** for any solution $\mathbf{f} = (f_1, f_2, \dots, f_N) \in \mathcal{F}_s$.


TSP: Finding the optimal tour with $\mathbf{f} \in \{0,1\}^N$ by selecting n edges from the total of $N = n^2$ edges

$$c_s^{\text{TSP}}(\mathbf{f}) = \sum_{i=1}^N f_i \cdot \text{length}(f_i, s)$$



MIS: Finding the maximal subset with $\mathbf{f} \in \{0,1\}^N$ by selecting a subset from the total of $N = n$ nodes.

$$c_s^{\text{MIS}}(\mathbf{f}) = \sum_{i=1}^N (1 - f_i)$$



4/2/2024

@Yiming Yang, 11-741 S24 Graph10 NAR CO Solvers

4

4

Objective

- We want to minimize the **expected cost of system-generated solutions** for graph s

$$\min_{\theta} \mathcal{L}_{\theta}(s) = \min_{\theta} \sum_{\pi \sim p_{\theta}(\cdot|s)} [c_s(\pi) p_{\theta}(\pi|s)] \quad (\text{previous lecture})$$

- Optimizing θ via reinforcement learning (RL) with

$$\nabla \mathcal{L}_{\theta}(s) = \sum_{\pi \sim p_{\theta}(\cdot|s)} c_s(\pi) \nabla \log p_{\theta}(\pi|s) \quad (\text{previous lecture})$$

- In AR model, $p_{\theta}(\cdot|s)$ is estimated by the decoder. **But now we want to eliminate the decode!**

- Idea:

- Training an **encoder-only** neural network to produce a **heatmap** (of node/edge weights based on training examples of $\langle \text{tour}, \text{cost} \rangle$ pairs).
- Using the heatmap to conduct a search of plausible solutions during testing.


Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

5

5

A toy example

Graph s Feasible Solutions & Costs ($n! = 4 \cdot 3 \cdot 2 \cdot 1 = 24$)

$$\pi_1 = 1 \rightarrow 2 \rightarrow 3 \rightarrow 4 \rightarrow 1, \quad c(\pi_1) = 200 + \textcolor{red}{200} + 400 + \textcolor{green}{400} = 1200$$

$$\pi_2 = 1 \rightarrow 2 \rightarrow 4 \rightarrow 3 \rightarrow 1, \quad c(\pi_2) = 200 + \textcolor{red}{300} + 400 + \textcolor{green}{100} = 1000$$

$$\pi_3 = \dots$$

$$\{(\pi_i, c_i)\} \xrightarrow{\text{intuitively}} \text{Is } e_{24} \text{ better than } e_{23}?$$

$$\{(\pi_i, c_i)\} \xrightarrow{\text{RL training}} \textcolor{red}{Heatmap} = \{\text{edge weights}\}$$

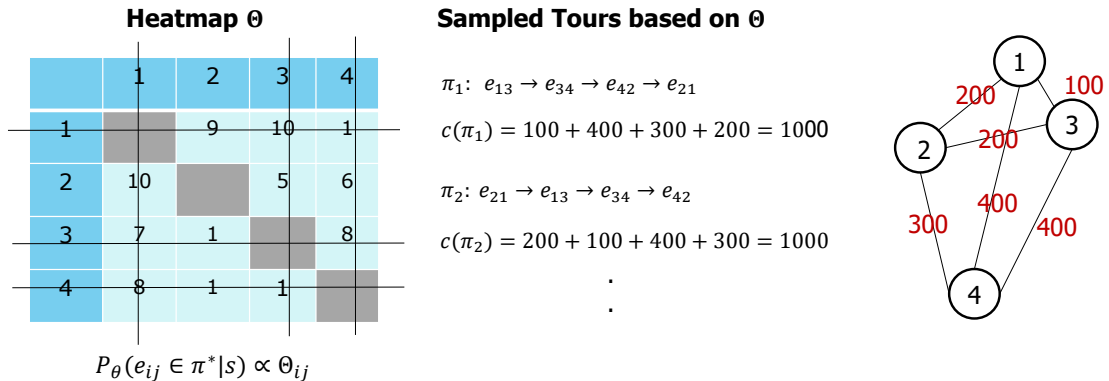
Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

6

6

Toy Example: NAR Sequence Generation in Testing Phase



Nov 14 & 19

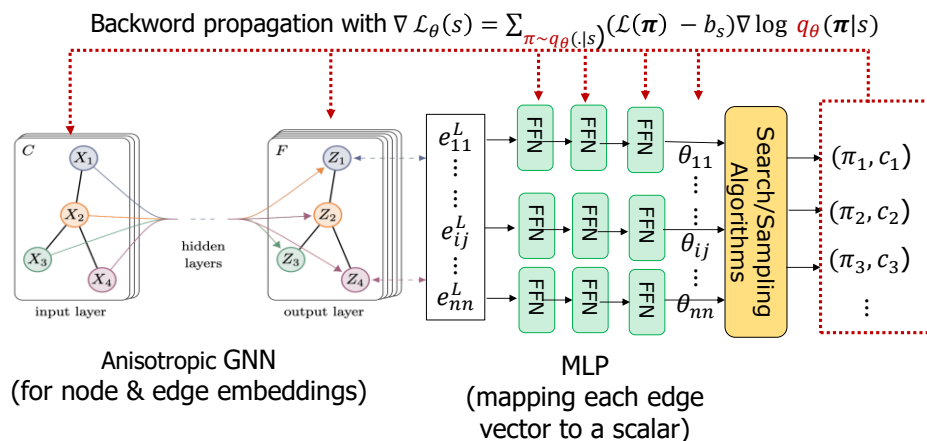
@Yiming Yang 11741 F24 Lecture on NAR CO Solver

7

7

How do we obtain the heatmap?

GNN+MLP with RL over a training set of graphs and sampled tours per graph



Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

8

8

Anisotropic Graph Neural Networks

producing edge embedding for solving TSP (Z Sun & Y Yang, NeurIPS 2022)

Anisotropic Graph Neural Networks We follow Joshi et al. [33] on the choice of neural architectures. The backbone of the graph neural network is an anisotropic GNN with an edge gating mechanism [9]. Let \mathbf{h}_i^ℓ and \mathbf{e}_{ij}^ℓ denote the node and edge features at layer ℓ associated with node i and edge ij , respectively. The features at the next layer is propagated with an anisotropic message passing scheme:

$$\mathbf{h}_i^{\ell+1} = \mathbf{h}_i^\ell + \alpha(\text{BN}(\mathbf{U}^\ell \mathbf{h}_i^\ell + \mathcal{A}_{j \in \mathcal{N}_i}(\sigma(\mathbf{e}_{ij}^\ell) \odot \mathbf{V}^\ell \mathbf{h}_j^\ell))), \quad (18)$$

$$\mathbf{e}_{ij}^{\ell+1} = \mathbf{e}_{ij}^\ell + \alpha(\text{BN}(\mathbf{P}^\ell \mathbf{e}_{ij}^\ell + \mathbf{Q}^\ell \mathbf{h}_i^\ell + \mathbf{R}^\ell \mathbf{h}_j^\ell)). \quad (19)$$

where $\mathbf{U}^\ell, \mathbf{V}^\ell, \mathbf{P}^\ell, \mathbf{Q}^\ell, \mathbf{R}^\ell \in \mathbb{R}^{d \times d}$ are the learnable parameters of layer ℓ , α denotes the activation function (we use SiLU [15] in this paper), BN denotes the Batch Normalization operator [30], \mathcal{A} denotes the aggregation function (we use mean pooling in this paper), σ is the sigmoid function, \odot is the Hadamard product, and \mathcal{N}_i denotes the outlinks (neighborhood) of node i . We use a 12-layer GNN with width 32.

Model Training

Loss Function

$$\mathcal{L}_\Phi(\mathcal{C}) = \sum_{s \in \mathcal{C}} l_q(\theta_s) = \sum_{s \in \mathcal{C}} l_q[F_\Phi(X_s, A_s)]$$

where \mathcal{C} is the training set of graphs and $s \in \mathcal{C}$ is an instance graph;

Φ is the neural network (GNN+MLP) parameter set;

$\theta_s := F_\Phi(X_s, A_s)$ is the heatmap produced by the neural network given input graph s ;

X_s is the matrix of node features and A_s is the adjacency matrix of graph s ;

$q = q(\theta_s)$ is the auxiliary distribution for efficient sampling of solutions from heatmap θ_s .

Gradient

$$\nabla_\Phi \mathcal{L}_\Phi(\mathcal{C}) = E_{s \in \mathcal{C}} [\nabla_\Phi \theta_s \cdot \nabla_{\theta_s} l_q(\theta_s)] = E_{s \in \mathcal{C}} [\nabla_\Phi F_\Phi(X_s, A_s) \cdot \nabla_{\theta_s} l_q(\theta_s)]$$

- **Forward/backward propagations:** updating Φ with the full formula and θ with the green part.

Results

Method	Type	TSP-500			TSP-1000			TSP-10000		
		Length ↓	Drop ↓	Time ↓	Length ↓	Drop ↓	Time ↓	Length ↓	Drop ↓	Time ↓
Concorde	OR (exact)	16.55*	—	37.66m	23.12*	—	6.65h	N/A	N/A	N/A
Gurobi	OR (exact)	16.55	0.00%	45.63h	N/A	N/A	N/A	N/A	N/A	N/A
LKH-3 (default)	OR	16.55	0.00%	46.28m	23.12	0.00%	2.57h	71.77*	—	8.8h
LKH-3 (less trails)	OR	16.55	0.00%	3.03m	23.12	0.00%	7.73m	71.79	—	51.27m
Nearest Insertion	OR	20.62	24.59%	0s	28.96	25.26%	0s	90.51	26.11%	6s
Random Insertion	OR	18.57	12.21%	0s	26.12	12.98%	0s	81.85	14.04%	4s
Farthest Insertion	OR	18.30	10.57%	0s	25.72	11.25%	0s	80.59	12.29%	6s
EAN	RL+S	28.63	73.03%	20.18m	50.30	117.59%	37.07m	N/A	N/A	N/A
EAN	RL+S+2-OPT	23.75	43.57%	57.76m	47.73	106.46%	5.39h	N/A	N/A	N/A
AM	RL+S	22.64	36.84%	15.64m	42.80	85.15%	63.97m	431.58	501.27%	12.63m
AM	RL+G	20.02	20.99%	1.51m	31.15	34.75%	3.18m	141.68	97.39%	5.99m
AM	RL+BS	19.53	18.03%	21.99m	29.90	29.23%	1.64h	129.40	80.28%	1.81h
GCN	SL+G	29.72	79.61%	6.67m	48.62	110.29%	28.52m	N/A	N/A	N/A
GCN	SL+BS	30.37	83.55%	38.02m	51.26	121.73%	51.67m	N/A	N/A	N/A
POMO+EAS-Emb	RL+AS	19.24	16.25%	12.80h	N/A	N/A	N/A	N/A	N/A	N/A
POMO+EAS-Lay	RL+AS	19.35	16.92%	16.19h	N/A	N/A	N/A	N/A	N/A	N/A
POMO+EAS-Tab	RL+AS	24.54	48.22%	11.61h	49.56	114.36%	63.45h	N/A	N/A	N/A
Att-GCN	SL+MCTS	16.97	2.54%	2.20m	23.86	3.22%	4.10m	74.93	4.39%	21.49m
DIMES (ours)	RL+G	18.93	14.38%	0.97m	26.58	14.97%	2.08m	86.44	20.44%	4.65m
	RL+AS+G	17.81	7.61%	2.10h	24.91	7.74%	4.49h	80.45	12.09%	3.07h
	RL+S	18.84	13.84%	1.06m	26.36	14.01%	2.38m	85.75	19.48%	4.80m
	RL+AS+S	17.80	7.55%	2.11h	24.89	7.70%	4.53h	80.42	12.05%	3.12h
	RL+MCTS	16.87	1.93%	2.92m	23.73	2.64%	6.87m	74.63	3.98%	29.83m
	RL+AS+MCTS	16.84	1.76%	2.15h	23.69	2.46%	4.62h	74.06	3.19%	3.57h

OR: Traditional Solver

RL: Reinforce. Learners

SL: Supervised Learners

Sampling Strategies:

G: Greedy Readout

S: Sampling

BS: Beam Search

AS: Active Search

MCTS: Monte-Carlo Tree Search

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

15

15

Outline of 3 Lectures: Graph 9, 10 and 11

- Introduction to ML for Combinatorial Optimization (CO)
- Autoregressive (AR) CO Solvers
- Non-autoregressive (NAR) CO Solvers
 - Reinforcement learning with DIMES [R Qiu*, Z Sun* & Y Yang, **NearIPS 2022**]
 - Supervised learning with DIFUSCO [Z Sun & Y Yang, **NearIPS 2023**]
- Pre-trained Large Language Models [C Yang et al., ICLR 2024]

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

16

16

DIFUSCO: Graph-based Diffusion Solvers for CO

[Z Sun & Y Yang, NeurIPS 2023]

- Borrowing the key idea from recent diffusion models for computer vision
- Developed the 1st diffusion-based solver for CO problems
- Beat other SOTA CO solvers (e.g., DIMES) in both accuracy and efficiency

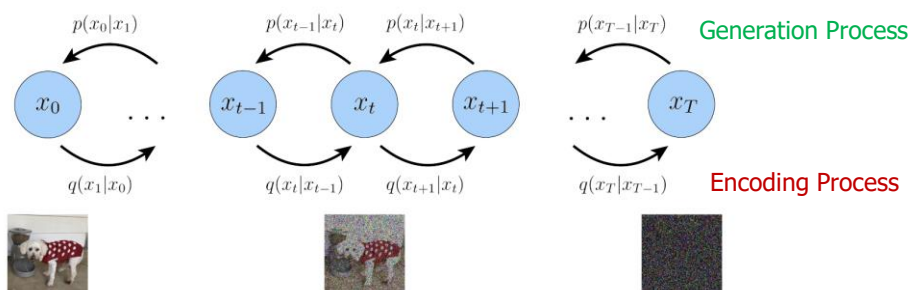
Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

17

17

Diffusion modeling for generating high-quality images



Rationale: Decomposing a hard problem ($x_T \rightarrow x_0$) into many easier problems ($x_t \rightarrow x_{t-1}$)

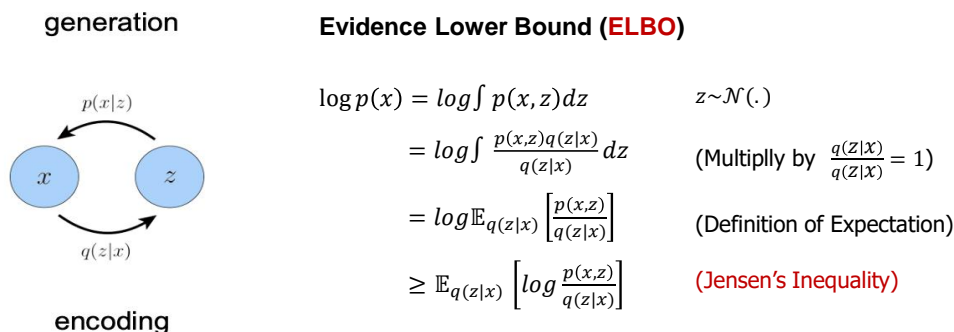
Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

18

18

Recap: Variational Autoencoder (VAE)



Maximization of $\log p(x)$ can be reduced to the problem of maximizing its ELBO.

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

19

19



Jensen's Inequality


- Jensen's inequality is a fundamental result in mathematics, particularly in convex analysis and probability theory. It states:

- If $f(x)$ is a **convex function** and X is a random variable, then

$$f(E[X]) \leq E[f(X)];$$

- If $f(x)$ is a **concave function**, then

$$f(E[X]) \geq E[f(X)].$$




This graph illustrates **Jensen's inequality** for $f(x) = x^2$ in the domain $[0, 3]$:

- The **blue curve** represents the convex function $f(x) = x^2$.
- The **orange dashed line** is the secant line connecting $(0, f(0))$ and $(3, f(3))$.

20

20

Diffusion Model as Markovian Hierarchical VAE (MHVAE)



Generation:
$$p(\mathbf{x}, \mathbf{z}_{1:T}) = p(\mathbf{z}_T) p_{\theta}(\mathbf{x} | \mathbf{z}_1) \prod_{t=2}^T p_{\theta}(\mathbf{z}_{t-1} | \mathbf{z}_t)$$

Encoding:
$$q_{\phi}(\mathbf{z}_{1:T} | \mathbf{x}) = q_{\phi}(\mathbf{z}_1 | \mathbf{x}) \prod_{t=2}^T q_{\phi}(\mathbf{z}_t | \mathbf{z}_{t-1})$$

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

21

21

ELBO of Diffusion Model (replacing \mathbf{z} by $\mathbf{z}_{1:T}$)

$$\begin{aligned} \log p(\mathbf{x}) &= \log \int p(\mathbf{x}, \mathbf{z}_{1:T}) d\mathbf{z}_{1:T} \\ &= \log \int \frac{p(\mathbf{x}, \mathbf{z}_{1:T}) q(\mathbf{z}_{1:T} | \mathbf{x})}{q(\mathbf{z}_{1:T} | \mathbf{x})} d\mathbf{z}_{1:T} && \text{(Multiply by } \frac{q(\mathbf{z} | \mathbf{x})}{q(\mathbf{z} | \mathbf{x})} = 1 \text{)} \\ &= \log \mathbb{E}_{q(\mathbf{z}_{1:T} | \mathbf{x})} \left[\frac{p(\mathbf{x}, \mathbf{z}_{1:T})}{q(\mathbf{z}_{1:T} | \mathbf{x})} \right] && \text{(Definition of Expectation)} \\ &\geq \mathbb{E}_{q(\mathbf{z}_{1:T} | \mathbf{x})} \left[\log \frac{p(\mathbf{x}, \mathbf{z}_{1:T})}{q(\mathbf{z}_{1:T} | \mathbf{x})} \right] && \text{(Apply Jensen's Inequality)} \end{aligned}$$

From now on, we change notation as $\mathbf{x} \rightarrow \mathbf{x}_0$, $\mathbf{z}_i \rightarrow \mathbf{x}_i$ and $p(\mathbf{x}, \mathbf{z}_{1:T}) \rightarrow p_{\theta}(\mathbf{x}_{0:T})$.

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

22

22

ELBO of Diffusion Model (notation changed)

$$\begin{aligned}
 \log p(\mathbf{x}) &= \log \int p(\mathbf{x}_{0:T}) d\mathbf{x}_{1:T} \\
 &= \log \int \frac{p(\mathbf{x}_{0:T}) q(\mathbf{x}_{1:T} | \mathbf{x}_0)}{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} d\mathbf{x}_{1:T} \\
 &= \log \mathbb{E}_{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \left[\frac{p(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \right] \\
 &\geq \underbrace{\mathbb{E}_{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \right]}_{\text{ELBO}}
 \end{aligned}$$

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver


23

23

ELBO of Diffusion Model

[Calvin Luo, 2022]

$$\begin{aligned}
 \log p(\mathbf{x}) &\geq \mathbb{E}_{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \right] \\
 &= \underbrace{\mathbb{E}_{q(\mathbf{x}_1 | \mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0 | \mathbf{x}_1)]}_{\text{reconstruction term}} - \underbrace{\mathcal{D}_{\text{KL}}(q(\mathbf{x}_T | \mathbf{x}_0) || p(\mathbf{x}_T))}_{\text{prior matching term}} \\
 &\quad - \underbrace{\sum_{t=2}^T \mathbb{E}_{q(\mathbf{x}_t | \mathbf{x}_0)} [\mathcal{D}_{\text{KL}}(q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) || p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t))]}_{\text{denoising matching term}}
 \end{aligned}$$



(Proof skipped: see the Appendix and the tutorial by Clavin Luo for details)

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

24

24

DIFUSCO as a generic (supervised) CO Solver

- We define a CO problem (e.g., TSP or MIS) as the task of generating discrete $\{0,1\}$ -vector $x \in \{0,1\}^{N(s)}$, which indicates a sequence (or set) of selected edges or nodes from an input graph, and $N(s)$ is the number of nodes or edges in the input graph.
- We train a neural network to minimize

$$L(\theta) = \mathbb{E}_{s \in S} [-\log p_{\theta}(x_s^*)]$$

where x_s^* is an optimal solution for graph s ; $p_{\theta}(x_s)$ is the system-estimated probability of any $x \in \{0,1\}^{N(s)}$ to be an optimal solution in this graph.

- We obtain a **labeled** training set for supervised learning of the model, with an existing heuristic (non-ML) solver that finds an (near-)optimal solution for each instance graph.

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

25

25

Forward Encoding (Discrete Model) (a pre-designed procedure, no learning)


- Manually specify hyperparameters $(\beta_1, \beta_2, \dots, \beta_T)$, namely the *schedule*, which contains the step-wise corruption ratios.
- Starting with clean data x_0 , inject noise in step $t = 1, \dots, T$ with for each element of x_t as

$$q(x_t^{(i)} | x_{t-1}^{(i)}) = \text{Cat}(x_t^{(i)} | p_t^{(i)} = \tilde{x}_{t-1}^{(i)} Q_t)$$

- Multi-step acceleration (e.g., randomly sample 50 time stamps out of $1, \dots, T=1000$)

$$q(x_t^{(i)} | x_{t-1}^{(i)}, x_0^{(i)}) = \text{Cat}(x_t^{(i)} | p_t^{(i)} = \tilde{x}_0^{(i)} \bar{Q}_t)$$

$$\bar{Q}_t \triangleq Q_1 Q_2 \dots Q_t$$



Efficient computation of x_t given x_0
(with closed-form instead of MCMC)

4/2/2024

@Yiming Yang, 11-741 S24 Graph10 NAR CO Solvers

26

26

Forward Encoding Details (Discrete Model)

- Each step, the element-wise corruption is defined as

$$q(x_t^{(i)} | x_{t-1}^{(i)}) = \text{Cat}(x_t^{(i)} | p_t^{(i)} = \tilde{x}_{t-1}^{(i)} Q_t)$$

- Applying it recursively yields

$$\begin{aligned} q(x_t^{(i)} | x_{t-1}^{(i)}, x_0^{(i)}) &= \text{Cat}(x_t^{(i)} | p_t^{(i)} = \tilde{x}_{t-1}^{(i)} Q_t) \\ &= \text{Cat}(x_t^{(i)} | p_t^{(i)} = \tilde{x}_{t-2}^{(i)} Q_{t-1} Q_t) \\ &= \text{Cat}(x_t^{(i)} | p_t^{(i)} = \tilde{x}_{t-3}^{(i)} Q_{t-2} Q_{t-1} Q_t) \\ &\quad \vdots \\ &= \text{Cat}(x_t^{(i)} | p_t^{(i)} = \tilde{x}_0^{(i)} Q_1 Q_2 \cdots Q_t) \\ &= \text{Cat}(x_t^{(i)} | p_t^{(i)} = \tilde{x}_0^{(i)} \bar{Q}_t) \end{aligned}$$

$$\bar{Q}_t \triangleq Q_1 Q_2 \cdots Q_t$$

$$\begin{aligned} p_t^{(i)} &= \begin{pmatrix} \tilde{x}_{t-1}^{(i)} & 1 - \tilde{x}_{t-1}^{(i)} \end{pmatrix} \begin{pmatrix} 1 - \beta_t & \beta_t \\ \beta_t & 1 - \beta_t \end{pmatrix} \\ &= \begin{cases} (1 - \beta_t, \beta_t) & \text{if } x_{t-1}^{(i)} = 1 \\ (\beta_t, 1 - \beta_t) & \text{if } x_{t-1}^{(i)} = 0 \end{cases} \end{aligned}$$

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

27


27

Backward Denoising (Generation)

(supervised learning with a **graph neural network**)

- Start with a fully noisy x_T whose elements $x_T^{(i)} \sim \text{Ber}(0.5)$ are sampled independently.
- For $t = T, \dots, 1$, calculate

$$p_\theta(x_{t-1} | x_t) = \sum_{\hat{x}_0} q(x_{t-1} | x_t, \hat{x}_0) p_\theta(\hat{x}_0 | x_t)$$



where $p_\theta(\hat{x}_0 | x_t)$ is estimated by a trained neural network for **multi-label classification** on labeled set $\mathcal{D} = \{(x_t, x_0)\}_{t=1}^{|D|}$, which takes x_t as the input and predicts $\hat{x}_0 \in \{0,1\}^N$ with $p_\theta(\hat{x}_0 | x_t)$.

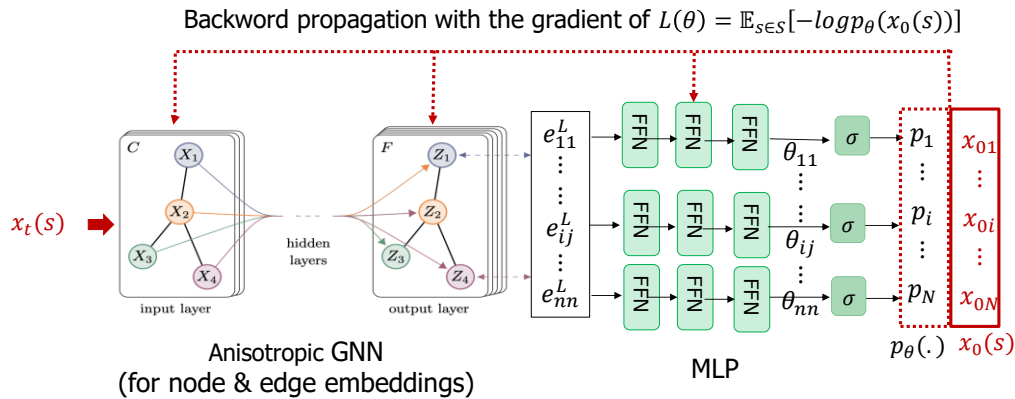
Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

30

30

Graph-based Denoising Network for $x_t(s) \xrightarrow{\text{mapping}} x_0(s)$



Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

31

31

Training the Denoising Network (details)

- Randomly sample K temporal stamps out of the T steps (e.g., $K=50$, $T=1000$);
- Compute x_t given x_0 for $t = 1$ to K using a close-form solution (slide #26);
- Train the GNN+MLP model (last slide) on labeled set $\mathcal{D} = \{(x_t, x_0)\}_{t=1}^K$;
- GNN input layers
 - For **TSP**, edge e_{ij}^0 is initialized as the corresponding value in $x_t(s)$, and node h_i^0 is initialized as sinusoidal features (e.g., the geographical coordinates) of the nodes in the input graph;
 - For **MIS**, edge e_{ij}^0 is initialized as zeros, and node h_i^0 is initialized as the corresponding values in $x_t(s)$;

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

32

32

Solution generation per graph in the testing phase

- Start with $x_T \in \{0,1\}^N$ with randomly sampled elements $x_T^{(i)} \sim \text{Ber}(0.5)$.
- Use the trained denoising network (GNN) to estimate

$$p_\theta(x_{t-1}^{(i)} | x_t^{(i)}) = \sum_{\hat{x}_0^{(i)} \in \{0,1\}} q(x_{t-1}^{(i)} | x_t^{(i)}, \hat{x}_0^{(i)}) p_\theta(\hat{x}_0^{(i)} | x_t^{(i)})$$
 - $q(x_{t-1}^{(i)} | x_t^{(i)}, \hat{x}_0^{(i)})$ is computed efficiently in a closed form (slide #26);
 - $p_\theta(\hat{x}_0^{(i)} | x_t^{(i)})$ is the output at the current step by the GNN.
- Sample $x_{t-1}^{(i)} \sim p_\theta(x_{t-1}^{(i)} | x_t^{(i)})$ for $i = 1, \dots, N$; set $t := t - 1$, and repeat above process until $t = 0$.
- Fast inference with multi-step acceleration can be done (e.g., collapsing every 50 steps into one).

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver







33

33

Demo of TSP Solution Generation

(Sun & Yang, NeurIPS 2023)

- Diffusion Model can generate diverse solutions



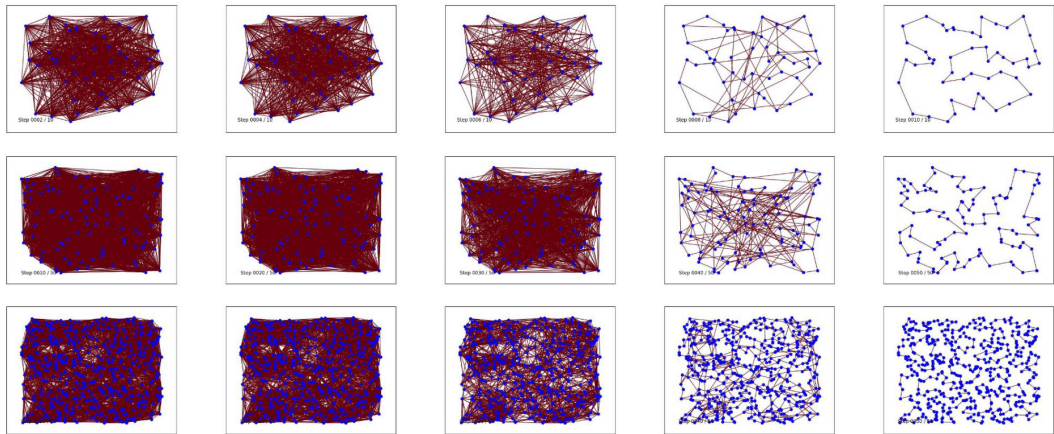
Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

35

35

DIFUSCO on TSP-50, TSP-100, TSP-500



Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

36

36

Results on Small Problems [Sun & Yang, NeurIPS 2023]

ALGORITHM	TYPE	TSP-50		TSP-100	
		LENGTH↓	GAP(%)↓	LENGTH↓	GAP(%)↓
CONCORDE*	EXACT	5.69	0.00	7.76	0.00
2-OPT	HEURISTICS	5.86	2.95	8.03	3.54
AM	GREEDY	5.80	1.76	8.12	4.53
GCN	GREEDY	5.87	3.10	8.41	8.38
TRANSFORMER	GREEDY	5.71	0.31	7.88	1.42
POMO	GREEDY	5.73	0.64	7.84	1.07
SYM-NCO	GREEDY	-	-	7.84	0.94
DPDP	1k-IMPROVEMENTS	5.70	0.14	7.89	1.62
IMAGE DIFFUSION	GREEDY [†]	5.76	1.23	7.92	2.11
OURS	GREEDY[†]	5.70	0.10	7.78	0.24
AM	1k×SAMPLING	5.73	0.52	7.94	2.26
GCN	2k×SAMPLING	5.70	0.01	7.87	1.39
TRANSFORMER	2k×SAMPLING	5.69	0.00	7.76	0.39
POMO	8×AUGMENT	5.69	0.03	7.77	0.14
SYM-NCO	100×SAMPLING	-	-	7.79	0.39
MDAM	50×SAMPLING	5.70	0.03	7.79	0.38
DPDP	100k-IMPROVEMENTS	5.70	0.00	7.77	0.00
OURS	16×SAMPLING	5.69	-0.01	7.76	-0.01

*Concorde is a TSP solver for integer coordinates.

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

37

37

Results on Large Problems [Sun & Yang, NeurIPS 2023]

ALGORITHM	TYPE	TSP-500			TSP-1000			TSP-10000		
		LENGTH ↓	GAP ↓	TIME ↓	LENGTH ↓	GAP ↓	TIME ↓	LENGTH ↓	GAP ↓	TIME ↓
CONCORDE	EXACT	16.55*	—	37.66m	23.12*	—	6.65h	N/A	N/A	N/A
GUROBI	EXACT	16.55	0.00%	45.63h	N/A	N/A	N/A	N/A	N/A	N/A
LKH-3 (DEFAULT)	HEURISTICS	16.55	0.00%	46.28m	23.12	0.00%	2.57h	71.77*	—	8.8h
LKH-3 (LESS TRAILS)	HEURISTICS	16.55	0.00%	3.03m	23.12	0.00%	7.73m	71.79	—	51.27m
FARTHEST INSERTION	HEURISTICS	18.30	10.57%	0s	25.72	11.25%	0s	80.59	12.29%	6s
AM	RL+G	20.02	20.99%	1.51m	31.15	34.75%	3.18m	141.68	97.39%	5.99m
GCN	SL+G	29.72	79.61%	6.67m	48.62	110.29%	28.52m	N/A	N/A	N/A
POMO+EAS-EMB	RL+AS+G	19.24	16.25%	12.80h	N/A	N/A	N/A	N/A	N/A	N/A
POMO+EAS-TAB	RL+AS+G	24.54	48.22%	11.61h	49.56	114.36%	63.45h	N/A	N/A	N/A
DIMES	RL+G	18.93	14.38%	0.97m	26.58	14.97%	2.08m	86.44	20.44%	4.65m
DIMES	RL+AS+G	17.81	7.61%	2.10h	24.91	7.74%	4.49h	80.45	12.09%	3.07h
OURS (DIFUSCO)	SL+G†	18.35	10.85%	3.61m	26.14	13.06%	11.86m	98.15	36.75%	28.51m
OURS (DIFUSCO)	SL+G†+2-OPT	16.80	1.49%	3.65m	23.56	1.90%	12.06m	73.99	3.10%	35.38m
EAN	RL+S+2-OPT	23.75	43.57%	57.76m	47.73	106.46%	5.39h	N/A	N/A	N/A
AM	RL+BS	19.53	18.03%	21.99m	29.90	29.23%	1.64h	129.40	80.28%	1.81h
GCN	SL+BS	30.37	83.55%	38.02m	51.26	121.73%	51.67m	N/A	N/A	N/A
DIMES	RL+S	18.84	13.84%	1.06m	26.36	14.01%	2.38m	85.75	19.48%	4.80m
DIMES	RL+AS+S	17.80	7.55%	2.11h	24.89	7.70%	4.53h	80.42	12.05%	3.12h
OURS (DIFUSCO)	SL+S	17.23	4.08%	11.02m	25.19	8.95%	46.08m	95.52	33.09%	6.59h
OURS (DIFUSCO)	SL+S+2-OPT	16.65	0.57%	11.46m	23.45	1.43%	48.09m	73.89	2.95%	6.72h
ATT-GCN	SL+MCTS	16.97	2.54%	2.20m	23.86	3.22%	4.10m	74.93	4.39%	21.49m
DIMES	RL+MCTS	16.87	1.93%	2.92m	23.73	2.64%	6.87m	74.63	3.98%	29.83m
DIMES	RL+AS+MCTS	16.84	1.76%	2.15h	23.69	2.46%	4.62h	74.06	3.19%	3.57h
OURS (DIFUSCO)	SL+MCTS	16.63	0.46%	10.13m	23.39	1.17%	24.47m	73.62	2.58%	47.36m

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

38

38

Concluding Remarks

- ~~Bad News:~~ ML solvers cannot beat exact solvers for small graphs ($n \leq 100$ nodes) because exact solvers guarantees to find optimal solutions when they can scale.
- Good News:** ML solvers, especially with the advanced NAR neural solvers (e.g., DIMES and DIFUSCO), can find **near-optimal solutions** for large problems that traditional **exact solvers fail to scale up**.
- Good News:** The advanced NAR neural solvers have achieved the performance level close to the best **domain-specific heuristic solvers but without hand-craft heuristics**.
- Follow-up efforts:** training diffusion-based models in unsupervised setting (without knowing the best solution per graph); pretrained LLMs (next)

Nov 14 & 19

@Yiming Yang 11741 F24 Lecture on NAR CO Solver

39

39

References

- [NeurIPS 2022] R Qiu*, Z Sun*, Z. and Y Yang. [DIMES: A differentiable meta solver for combinatorial optimization problems](#)
- [NeurIPS 2023] Sun, Z. and Yang, Y. [Difusco: Graph-based Diffusion Solvers for Combinatorial Optimization](#)
- [Tutorial 2022] Calvin Luo. Understanding Diffusion Models: [A Unified Perspective: An intuitive, accessible tutorial on diffusion models.](#)

Appendix: ELBO of the Diffusion Model

$$\begin{aligned}
 \log p(\mathbf{x}) &\geq \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T) \prod_{t=1}^T p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{\prod_{t=1}^T q(\mathbf{x}_t|\mathbf{x}_{t-1})} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T) p_\theta(\mathbf{x}_0|\mathbf{x}_1) \prod_{t=2}^T p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_1|\mathbf{x}_0) \prod_{t=2}^T \boxed{q(\mathbf{x}_t|\mathbf{x}_{t-1})}} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T) p_\theta(\mathbf{x}_0|\mathbf{x}_1) \prod_{t=2}^T p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_1|\mathbf{x}_0) \prod_{t=2}^T \boxed{q(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0)}} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p_\theta(\mathbf{x}_T) p_\theta(\mathbf{x}_0|\mathbf{x}_1)}{q(\mathbf{x}_1|\mathbf{x}_0)} + \log \prod_{t=2}^T \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{\boxed{q(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0)}} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T) p_\theta(\mathbf{x}_0|\mathbf{x}_1)}{q(\mathbf{x}_1|\mathbf{x}_0)} + \log \prod_{t=2}^T \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{\boxed{\frac{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) q(\mathbf{x}_t|\mathbf{x}_0)}{q(\mathbf{x}_{t-1}|\mathbf{x}_0)}}} \right]
 \end{aligned}$$

Due to Markov property

Due to Bayes rule

Appendix: ELBO of the Diffusion Model

$$\begin{aligned}
& \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T)p_\theta(\mathbf{x}_0|\mathbf{x}_1)}{q(\mathbf{x}_1|\mathbf{x}_0)} + \log \prod_{t=2}^T \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{\frac{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)q(\mathbf{x}_t|\mathbf{x}_0)}{q(\mathbf{x}_{t-1}|\mathbf{x}_0)}} \right] \\
&= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T)p_\theta(\mathbf{x}_0|\mathbf{x}_1)}{q(\mathbf{x}_1|\mathbf{x}_0)} + \log \prod_{t=2}^T \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{\frac{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)q(\mathbf{x}_t|\mathbf{x}_0)}{q(\mathbf{x}_{t-1}|\mathbf{x}_0)}} \right] \\
&= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T)p_\theta(\mathbf{x}_0|\mathbf{x}_1)}{q(\mathbf{x}_1|\mathbf{x}_0)} + \log \frac{q(\mathbf{x}_1|\mathbf{x}_0)}{q(\mathbf{x}_T|\mathbf{x}_0)} + \log \prod_{t=2}^T \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\
&= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T)p_\theta(\mathbf{x}_0|\mathbf{x}_1)}{q(\mathbf{x}_T|\mathbf{x}_0)} + \sum_{t=2}^T \log \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\
&= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)] + \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T)}{q(\mathbf{x}_T|\mathbf{x}_0)} \right] + \sum_{t=2}^T \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\
&= \mathbb{E}_{q(\mathbf{x}_1|\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)] + \mathbb{E}_{q(\mathbf{x}_T|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T)}{q(\mathbf{x}_T|\mathbf{x}_0)} \right] + \sum_{t=2}^T \mathbb{E}_{q(\mathbf{x}_t, \mathbf{x}_{t-1}|\mathbf{x}_0)} \left[\log \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\
&= \underbrace{\mathbb{E}_{q(\mathbf{x}_1|\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)]}_{\text{reconstruction term}} - \underbrace{D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0) \parallel p(\mathbf{x}_T))}_{\text{prior matching term}} - \underbrace{\sum_{t=2}^T \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)} [D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t))]}_{\text{denoising matching term}}
\end{aligned}$$

4/2/2024

@Yiming Yang, 11-741 S24 Graph10 NAR CO Solvers

42