

# Neural Network Solvers for Combinatorial Optimization

---

## Graph 11. Other CO Solvers (LLMs)

---

1

## Lectures on Neural-network CO Solvers

---

- Graph 9: Autoregressive (AR) CO Solvers
- Graph 10: Non-autoregressive (NAR) CO Solvers
- Graph 11: Pre-trained Large Language Models for CO
- Graph 12: Neural Solvers for Mixed Integer Programming

2

## Large Language Models as Optimizers

(Google DeepMind: C Yang\*, C Chen\*, et al., ICLR 2024)

### Key Idea

- Proposing **OPRO** (Optimization by Prompting) as a generic optimizer for solving **any problems** (e.g., regression or TSP) described in natural language;
- Each prompt contains a task description and a few solution/score pairs for previously solved problem instances;
- Using a LLM (with the prompt) to generate solutions for each new problem instance;
- Evaluate each new solution and adding the new solution/score pair to the prompt;
- Repeating the above steps until a termination condition is met.

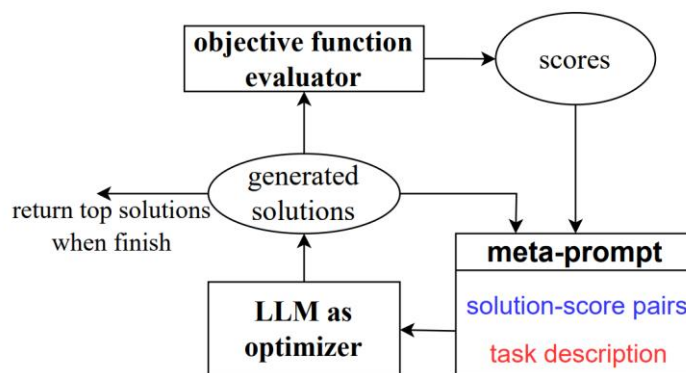
4/4/2024

@Yiming Yang, 11-741 S24 Graph11 Other CO Solvers (LLMs)

3

3

## Learn to Solve Problems with OPRO



Yang, Chengrun, et al. "Large language models as optimizers." arXiv preprint arXiv:2309.03409 (2023).

4

4

C.1 META-PROMPT FOR MATH OPTIMIZATION




Figure 19: An example of the meta-prompt for linear regression. The blue text contains solution-score pairs; the orange text are meta-instructions.

5

5

# Meta-prompt for TSP

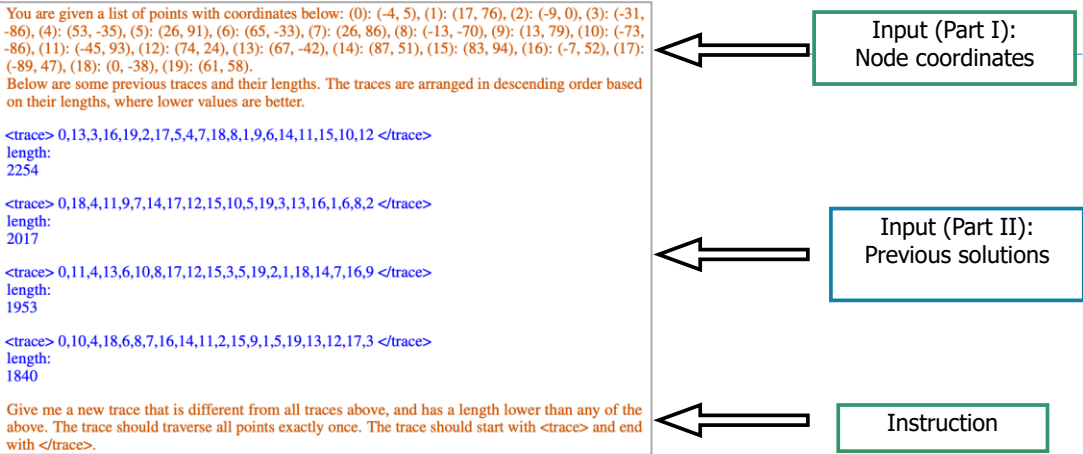


Figure 18: An example of the meta-prompt for Traveling Salesman Problems with problem size  $n = 20$ . The blue text contains solution-score pairs; the orange text are meta-instructions.

6

Yang, Chengrun, et al. "Large language models as optimizers." arXiv preprint arXiv:2309.03409 (2023).

6

## Evaluation Results on TSP

n	optimality gap (%)					# steps (# successes)		
	NN	FI	text-bison	gpt-3.5-turbo	gpt-4	text-bison	gpt-3.5-turbo	gpt-4
10	13.0 ± 1.3	3.2 ± 1.4	<b>0.0</b> ± 0.0	<b>0.0</b> ± 0.0	<b>0.0</b> ± 0.0	40.4 ± 5.6 (5)	46.8 ± 9.3 (5)	<b>9.6</b> ± 3.0 (5)
15	9.4 ± 3.7	1.2 ± 0.6	4.4 ± 1.3	1.2 ± 1.1	<b>0.2</b> ± 0.2	N/A (0)	202.0 ± 41.1 (4)	<b>58.5</b> ± 29.0 (4)
20	16.0 ± 3.9	<b>0.2</b> ± 0.1	30.4 ± 10.6	4.4 ± 2.5	1.4 ± 0.6	N/A (0)	438.0 ± 0.0 (1)	<b>195.5</b> ± 127.6 (2)
50	19.7 ± 3.1	<b>9.8</b> ± 1.5	219.8 ± 13.7	133.0 ± 6.8	11.0 ± 2.6	N/A (0)	N/A (0)	N/A (0)

- Baseline NN (Nearest Neighbor Heuristic)
  - At each step, select the closest node from the current partial solution
- Baseline FI (Farthest Insertion)
  - At each step, add a new node that maximize the minimal insertion cost which is defined as

$$c(k) = \min_{i,j} d(i, k) + d(k, j) - d(i, j)$$

4/4/2024

@Yiming Yang, 11-741 S24 Graph11 Other CO Solvers (LLMs)

7

7

## Concluding Remarks

- **Concept proving**
  - ORPO shows that LLMs with prompts in a loop can learn to optimize (mimicking gradient descent?)
- **Main limitations**
  - It cannot scale to large graphs or large training set of <solution, value> pairs.
- **Strong baselines are missing**
  - Comparison with DIMES and DIFUSCO on graphs with n=10000 nodes?
  - Comparison with classic exact solvers?
  - Comparison with LLMs for code generation?

4/4/2024

@Yiming Yang, 11-741 S24 Graph11 Other CO Solvers (LLMs)

8

8